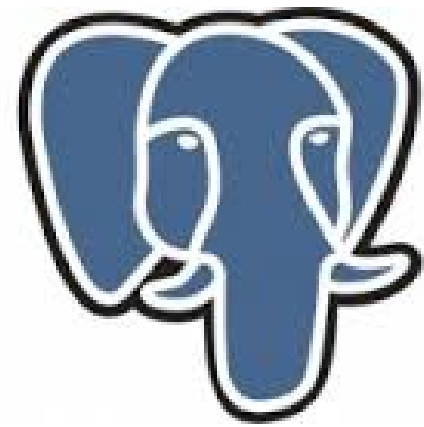


# Gimme some YESQL !



- and a GIS -



**PostgreSQL ?**  
**What's that ?**



**« Oh, yeah, this thing bundled with PostGIS  
installer ! »**



# **RODBMS**

**Relational  
Object  
Database Management System**

**SQL compliant  
Transactionnal**

**ACID**

**MVCC**

**Extensible**

**...**

**Spatial !**



# PostgreSQL

## New stuff

(..or some not that old features..)



# PG 9.x



# PostgreSQL

- › **PG 9.4 beta 2**
- › **Not long before release**
- › **Good foundation work**
- › **Lots of new features**
- › **Some of them useful for GIS**



# SQL powah !





# LATERAL JOIN

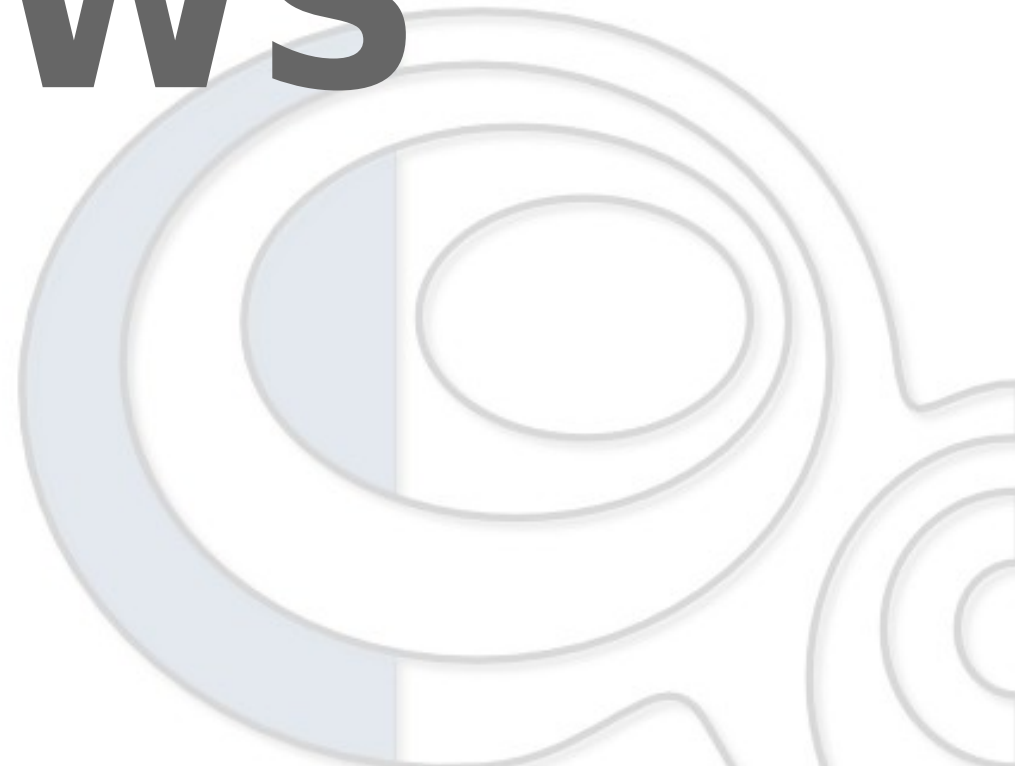


# LATERAL JOIN

- › Access to other tables in subqueries
- › Fix the < - > operator constant issue

```
-- find 2 bars closest to each bus stop (just in case one is closed)
select
    bus.gid, bus.name, lat.gid, lat.name, lat.dist
from
    points as bus
, lateral (
    select
        bar.gid
        , st_distance(bar.geom, bus.geom) as dist
        , bar.name
    from
        points as bar
    where
        bar.type = 'bar'
    order by
        bar.geom <-> bus.geom -- forbidden without lateral
    limit 2
) as lat
where
    bus.type = 'bus_stop'
order by
    bus.gid, lat.dist desc;
```

# **AUTO-UPDATE VIEWS**



# Auto-updateable views

- › **9.3+**
- › **No trigger !**
- › **9 .4 : column-level management**
  - › **Mix updateable/non updateable columns**
- › **Fine-tuning spatial right management**
- › **With check option**



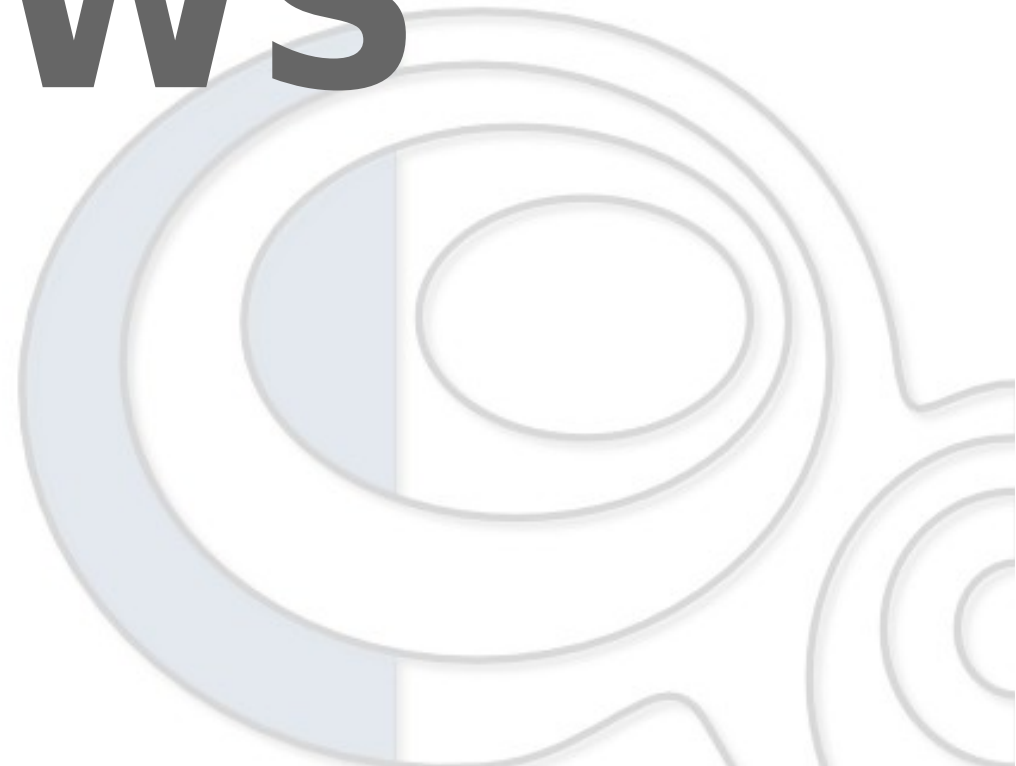
# Auto-updateable views

```
create or replace view
  ways_v as
select
  *
from
  ways
where
  st_contains(st_setsrid(st_geomfromtext(
    'polygon((4.8300 45.7562, 4.8300 45.7616, 4.8385 45.7616
    , 4.8385 45.7562, 4.8300 45.7562))'), 4326)
    , linestring)
with check option;


-- simplify geometry
update ways_v set linestring = st_simplify(linestring, 0.00001);

-- Check option will prevent this one
update ways_v set linestring = st_translate(linestring, -10, -10);
```

# **MATERIALIZED VIEWS**



# Materialized views

- › **9.3+**
  - › **Actually compute the view & store results**
  - › **« Cache views »**
  - › **Allows indexing**
  - › **Indexes & Constraints NOT copied**
  - › **9.4 : Refresh concurrently**
  - › **Use cases**
    - › **MV of complex joins**
    - › **MV of complex spatial operations**
    - › **MV of FDW tables**
- 

# Materialized views

```
-- initial table
create table pts (id serial, geom geometry(Point, 2154));
--fill it
insert into
    pts (geom)
select
    st_setsrid(st_makepoint(random() * 100000, random() * 100000), 2154)
from
    generate_series(1, 1000);
```

```
-- MV for long computation
create materialized view
    pts_buf as
select
    id
    , st_buffer(geom, 10) as geom
from
    pts;

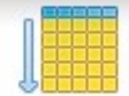
-- Data is updated
update pts set geom = st_translate(geom, 100, 100);

-- Refresh the MV|
refresh materialized view pts_buf;
```

## Seq Scan

on pts\_buf

(cost=0.00..87.00 rows=1000 width=572)



pts\_buf



# RANGES



# Ranges

- › **Datatype(s)**
- › **Integer, numeric, dates**
- › **Open, closed, infinite**
- › **Use indexes**
- › **Good for time data**

```
=# SELECT * from test_range ;
                        period
-----
["2012-01-01 00:00:00+01","2012-01-02 12:00:00+01"]
["2012-01-01 00:00:00+01","2012-03-01 00:00:00+01"]
["2008-01-01 00:00:00+01","2015-01-01 00:00:00+01"]
(3 rows)

=# SELECT * FROM test_range WHERE period && '[2012-01-03 00:00:00,2012-01-03 12:00:00]';
                        period
-----
["2012-01-01 00:00:00+01","2012-03-01 00:00:00+01"]
["2008-01-01 00:00:00+01","2015-01-01 00:00:00+01"]
(2 rows)
```

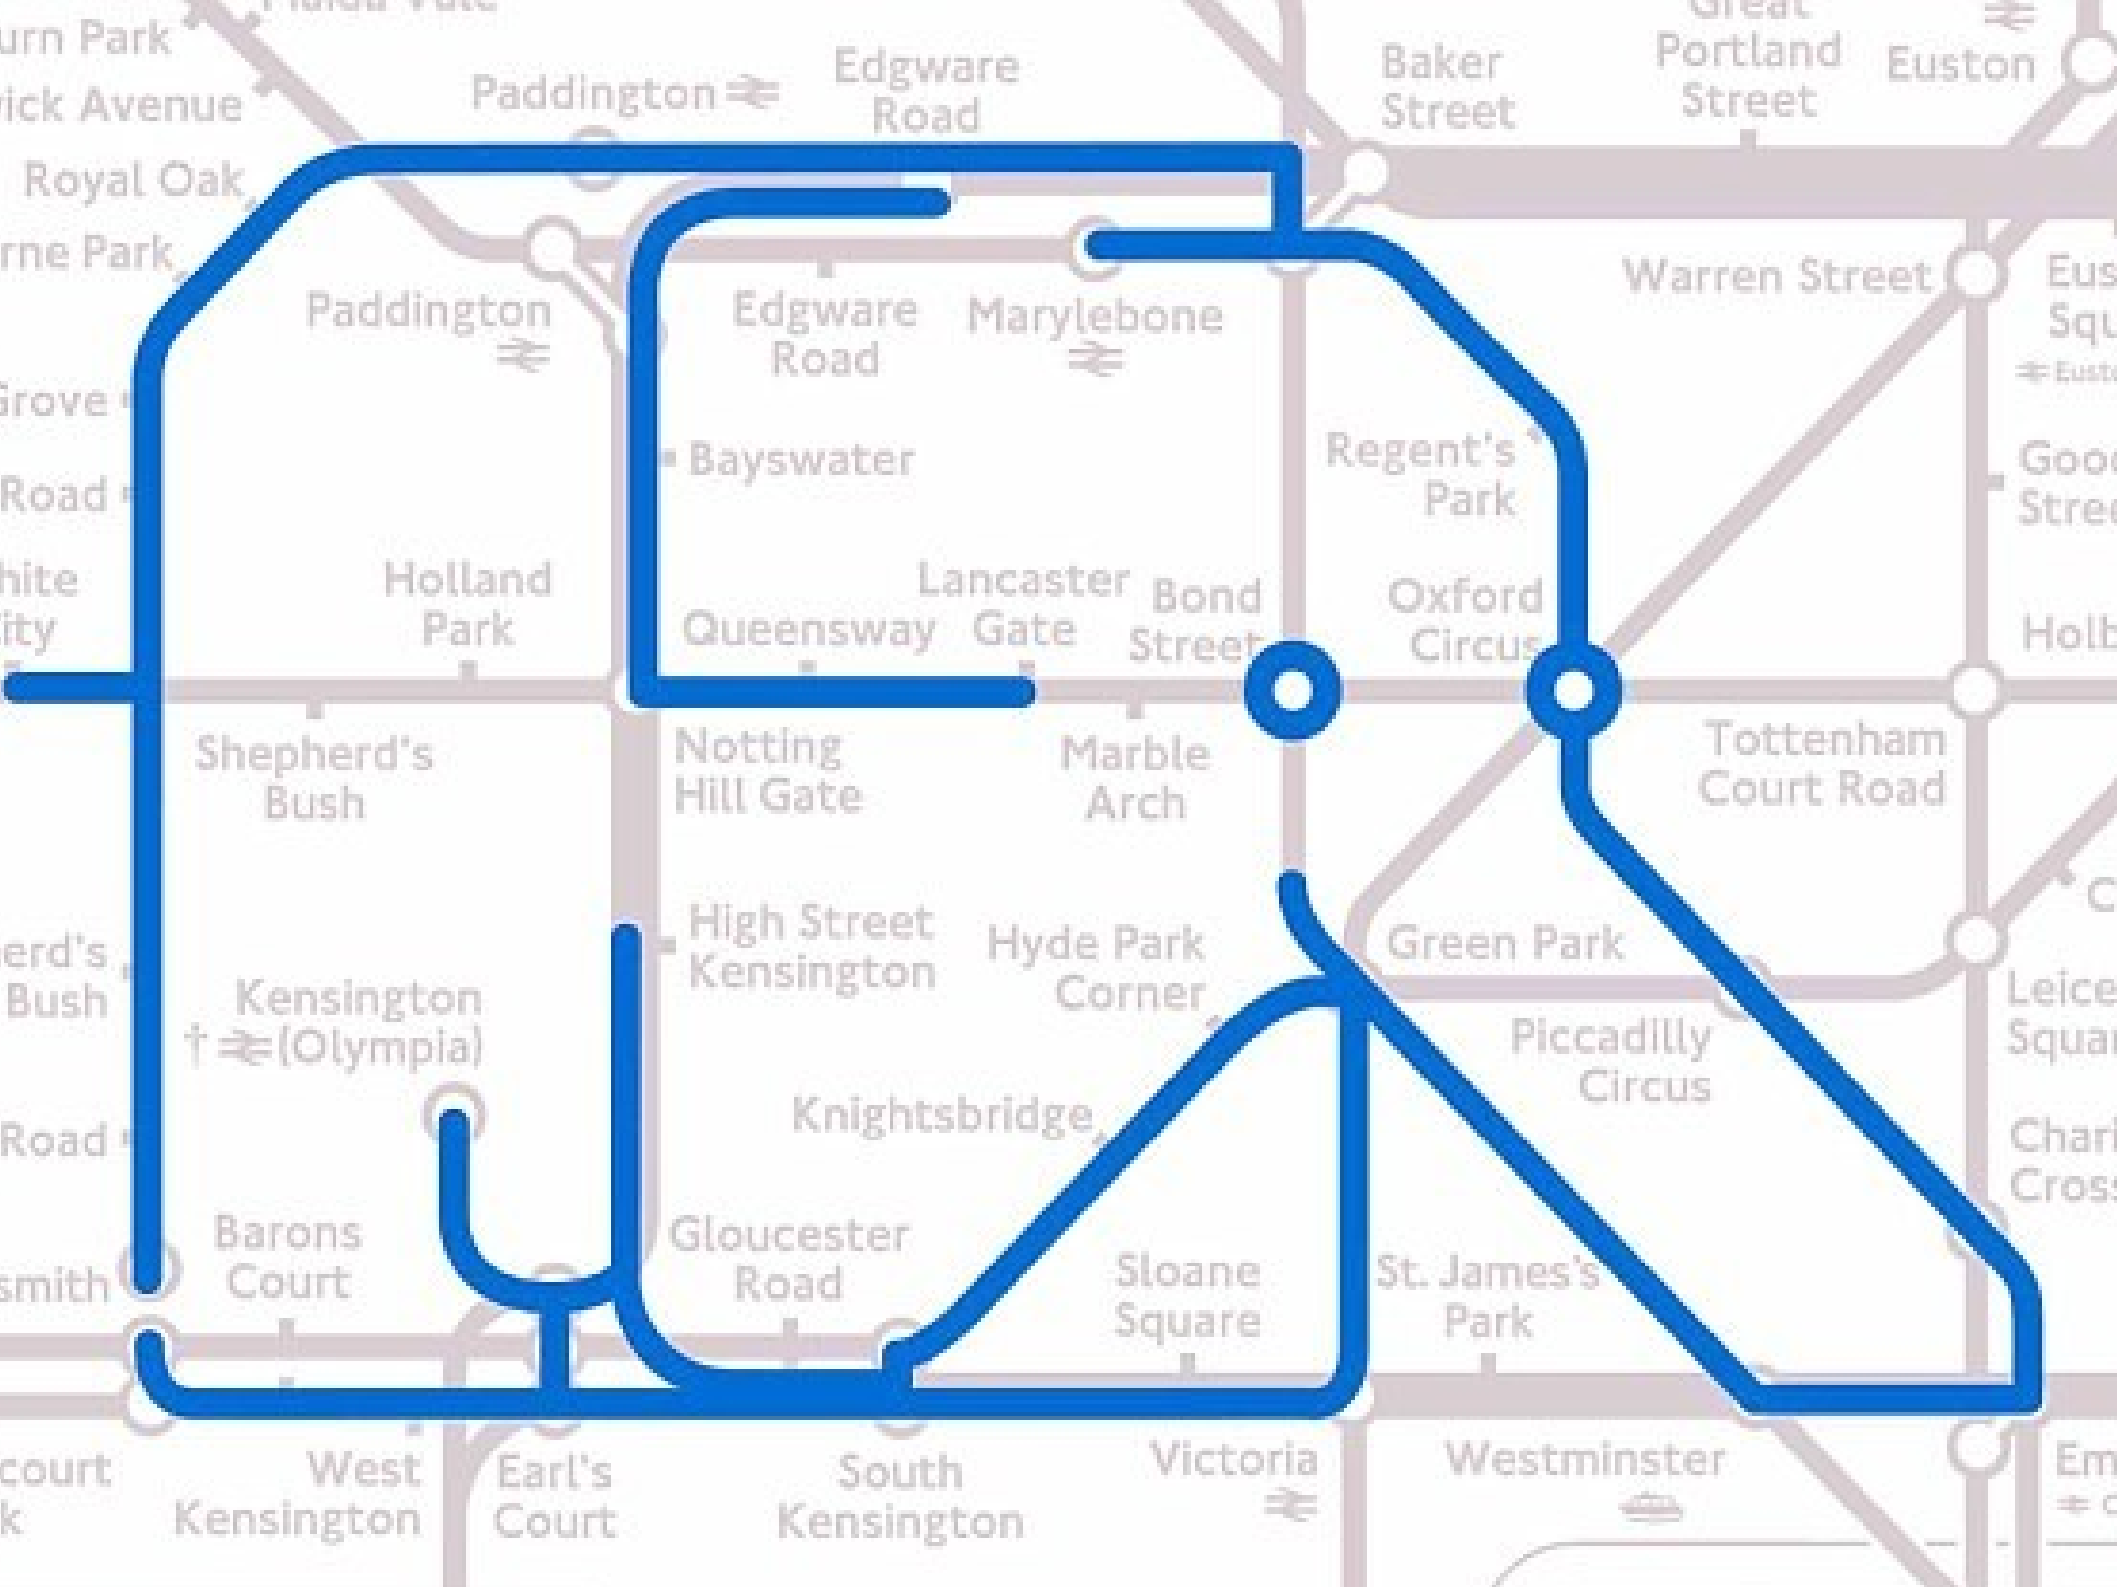
# WRITABLE CTE



```
CREATE TABLE bad_data (  
    id bigint geom geometry(MultiPolygon, 2154)  
);  
  
WITH deleted AS (  
    DELETE FROM  
        good_gis_data  
    WHERE  
        not st_isvalid(geom)  
    RETURNING  
        id, geom)  
INSERT INTO  
    bad_data  
SELECT  
    *  
FROM  
    deleted;
```

# RECURSIVE CTE





# Recursive CTEs

- › **Common Table Expression**
- › **SQL Standard**
- › **« WITH RECURSIVE »**
- › **CTE Syntax :-)**
- › **CTE Performances /!\**
- › **Recursive → Graph analysis features**

```
CREATE RECURSIVE VIEW t(n) AS  
  VALUES (1)  
  UNION ALL  
  SELECT n+1 FROM t WHERE n < 100;
```



```

create table
  rec_res as
with recursive
  search_graph(gid, source, depth, path, length, cycle) as (
    select
      g.gid, g.source, 1 as depth, ARRAY[g.gid] as path
      , cost, false as cycle
    from
      tr as g
    where
      gid = 31913
    union all
    select
      g.gid
      , g.source
      , sg.depth + 1 as depth
      , path || g.gid as path
      , sg.length + g.cost as length
      , g.gid = ANY(path) as cycle
    from
      tr as g
    join
      search_graph as sg
    on
      sg.source = g.target
    where
      not cycle
  )

```

1

2

## Recursive CTE

```

select
  sg.*
  , tr.geom
from
  search_graph as sg
join
  tr
on
  sg.gid = tr.gid
limit 1000;

```

3

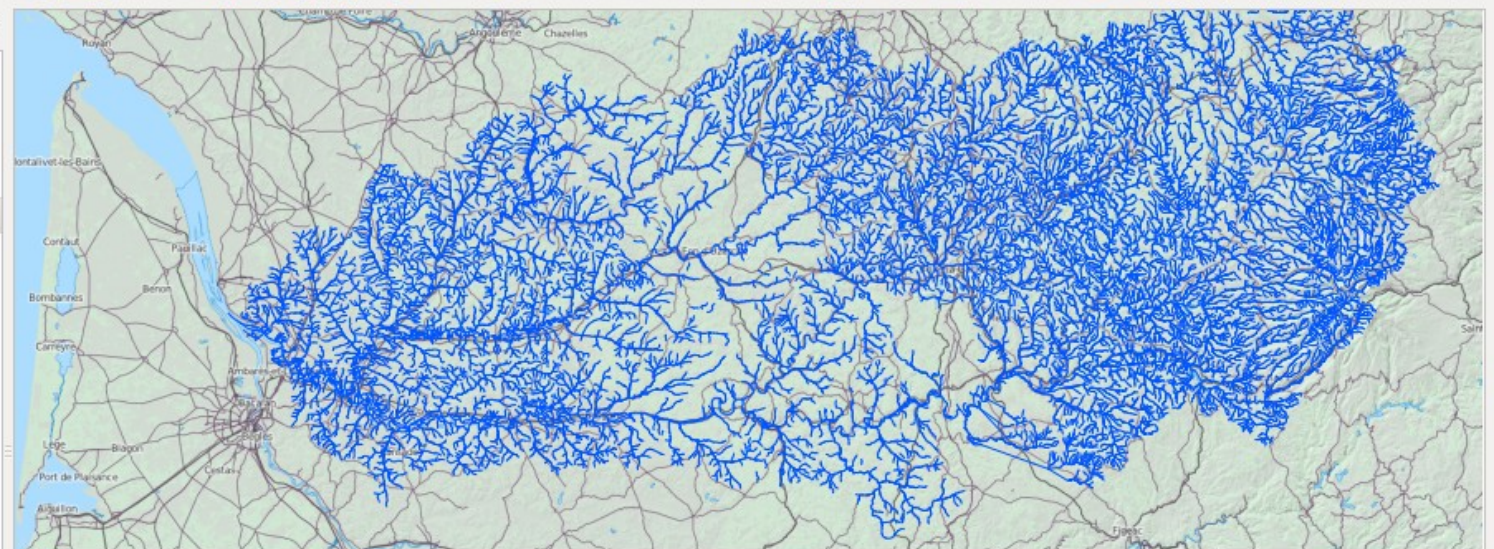
gid integer	source integer	depth integer	path integer[]	length double precision	cycle boolean	geom geometry(MultiLineString,2154)
31913	20850	1	{31913}	2666.0523017	f	01050000206A080000001000
33855	20735	2	{31913,	3473.3086319	f	01050000206A080000001000
32477	20845	2	{31913,	2725.7640259	f	01050000206A080000001000
33854	19909	3	{31913,	7183.7295195	f	01050000206A080000001000



Couches

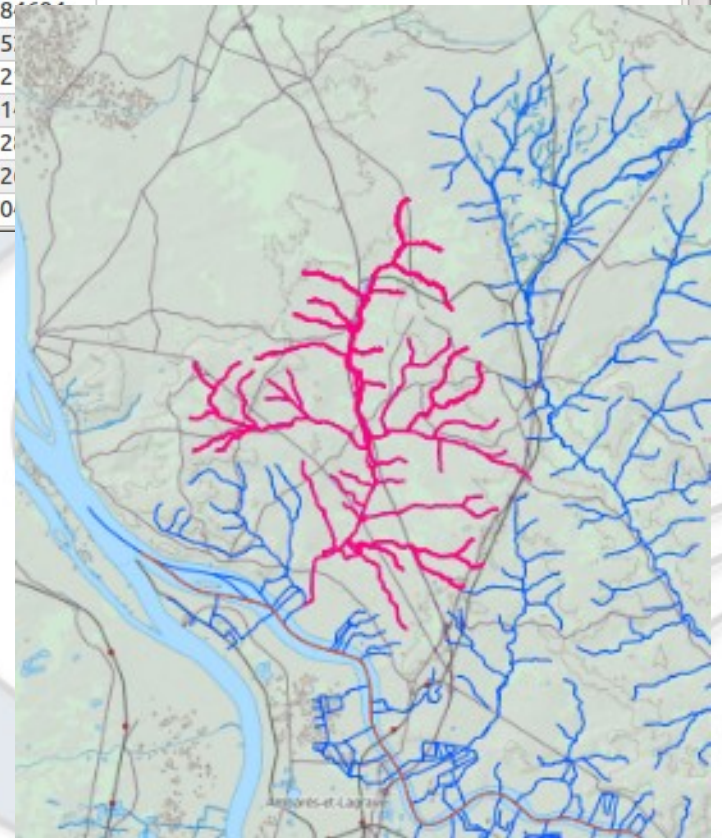
- ☐ recursive\_upstream\_topo
- ☐ recursive\_upstream
- ☐ shortest\_path\_topology
- ☐ shortest\_path\_pgouting
- ☒ hydro network
- ☒ background

☒ Contrôle de l'ordre de rendu des couches



Attribute table - hydro network :: 0 / 18936 feature(s) selected

	gid	source	target	hname	cost
0	17681	3042	3041	ruisseau de...	13.1468627...
1	50006	4363	4376	ruisseau de...	154.831357...
2	107308	4427	4443	ruisseau la ...	70.478...
3	110767	4810	4816	ruisseau le ...	426.45...
4	8923	4892	4827	ruisseau de...	1648.2...
5	109594	5158	5264	rivière la di...	946.01...
6	45039	5407	5429	NULL	114.02...
7	105937	5480	5594	ruisseau le ...	824.62...
8	104620	5481	5518	ruisseau la ...	243.00...



# UNLOGGED TABLES



# Unlogged tables

- › **No WAL**
- › **No maintenance**
- › **Faster to write to**
- › **will not survive a crash**
- › **will not survive a crash**
- › **WILL NOT SURVIVE A CRASH**
  
- › **Temporary results**
- › **Reproduceable results**



# Unlogged tables

```
-- Unlogged table  
CREATE UNLOGGED table  
    testu (a int, g geometry(LineString, 2154));  
  
-- 9.5 :  
alter table testu set logged;
```



# LOGICAL DECODING





# Logical Decoding

- › **Probably biggest 9.4 feature**
- › **First step of many**
- › **A kind of « database modification logger»**
- › **Open path to :**
  - › **Logical replication**
  - › **Multi-master replication**
  - › **Logging, versioning, incremental updates...**



# Logical Decoding

```
poggis=# SELECT 'init' FROM pg_create_logical_replication_slot('oslandia', 'test_decoding');
?column?
```

```
-----
init
(1 row)
```

```
poggis=# delete from points where gid = 1;
```

```
DELETE 1
```

```
poggis=# SELECT * FROM pg_logical_slot_get_changes('oslandia', null, null, 'include-xids', '0');
```

```
location | xid | data
-----+-----+-----
0/5A799658 | 931 | BEGIN
0/5A799658 | 931 | table public.points: DELETE: gid[integer]:1
0/5A79B6E8 | 931 | COMMIT
(3 rows)
```

```
poggis=# update points set geom = st_translate(geom, 1, 1) where gid = 2;
```

```
UPDATE 1
```

```
poggis=# SELECT * FROM pg_logical_slot_get_changes('oslandia', null, null, 'include-xids', '0');
```

```
poggis=#
location | xid |
-----+-----+
0/5A79E108 | 933 | BEGIN
0/5A79E108 | 933 | table public.points: INSERT: gid[integer]:999999999
      osm_id[character varying]:'99999999' "timestamp"[character varying]:''
      name[character varying]:'testrep'
      type[character varying]:'notype'
      geom[geometry]:'0101000020E610000000000000286A37B43F000000BD5B0BE53F'
0/5A7A0B08 | 933 | COMMIT
(3 rows)
```

# YESQL !





# **NoSQL is the future ? What NoSQL ?**

**→ PostgreSQL document-oriented data  
support**

**Unstructured data,  
key-pair values**

**« MongoDB on ACID »**

**~~XML~~, [hstore], JSON, JSONB**

# JSON



# JSON

- › 9.2+
- › Document storage
- › key order + duplicates preserved
- › Indexing
- › Text storage
- › I/O fast, size overhead
- › Operations < hstore
- › Expression index
- › **No GIN / GIST**

```
% CREATE INDEX idx_dvd_reviews  
ON reviews ((review#>>'{product,group}'));
```

# JSON

## > Operators & Functions (9.3, 9.4)

```
% SELECT ' [{"a":"foo"}, {"b":"bar"}, {"c":"baz"} ] ' :: json->2;  
?column?
```

```
-----  
{ "c": "baz" }
```

```
% SELECT ' { "a": { "b": "foo" } } ' :: json->'a';  
?column?
```

```
-----  
{ "b": "foo" }
```

```
% SELECT ' { "a": 1, "b": 2 } ' :: json->>'b';  
?column?
```

```
-----  
2
```

```
% SELECT ' [ { "a": { "b": { "c": "foo" } } } ] ' :: json#>' {a,b} ' ;  
?column?
```

```
-----  
{ "c": "foo" }
```

```
poggis=# select to_json(points) from points limit 1;  
{"gid":3,"osm_id":"124585","timestamp":"2014-04-04T08:17:23Z",  
"name":"Espace Comboire Nord","type":"motorway_junctio",  
"geom":"0101000020E61000009F2523C21ACA1640DE5DC2FC70934640"}
```

# JSON

## Other functions

- › **to\_json()**
- › **json\_build\_array()**
- › **json\_build\_object()**
- › **json\_each()**
- › **json\_each\_text()**
- › **json\_array\_length()**
- › **json\_object\_keys()**
- › **json\_array\_elements()**
- › **json\_array\_elements\_text()**
- › **json\_typeof()**
- › **json\_to\_record()**



# JSONB



# JSONB

- › 9.4+
- › Binary storage
- › No key duplicates & order preservation
- › Slower I/O, bigger than JSON
- › Hstore-like operators
- › JSON operators
- › .. and more : =, @>, ?, ?|, ?&
- › **GIN indexing**
- › Use it !
  - › Except I/O, key preservation (JSON)  
or flat + strings (hstore)



# JSONB

- › **GIN indexing** : @> ? ?| ?&
- › **Fast**
- › **Expression indexes**
- › **jsonb\_path\_ops** for better @>

```
CREATE INDEX idx_docs_gin  
ON docs USING gin(doc jsonb_path_ops);
```



# JSONB

```
% SELECT '{"a": 1, "b": 2}'::jsonb =  
           '{"b": 2, "a": 1}'::jsonb;  
?column?  
-----  
t  
% SELECT '{"a":1, "b":2}'::jsonb @>  
           '{"b":2}'::jsonb;  
?column?  
-----  
t  
% SELECT '{"a":1, "b":2}'::jsonb ? 'b';  
?column?  
-----  
t  
%  
% SELECT '{"a": {"b": 2, "c": 3}}'::jsonb @>  
           '{"a": {"c": 3}}'::jsonb;  
?column?  
-----  
.  
% SELECT '{"a":1, "b":2, "c":3}'::jsonb ?| ARRAY['b', 'd'];  
?column?  
-----  
t
```

# GeoJSON

› **PostGIS plays too !**

```
select st_asgeojson(geom)::json from points limit 1;
```

.....

```
{ "type": "Point", "coordinates": [5.6973677, 45.1518856] }  
(1 row)
```

```
select  
  row_to_json(f) as feature  
from (   
  select  
    'Feature' as type  
    , st_asgeojson(geom)::json as geometry  
    , '{"gid":1, "color" : "green"}'::json as properties  
  from  
    points  
  limit 1  
) as f;
```

```
{  
  "type": "Feature",  
  "geometry": {  
    "type": "Point",  
    "coordinates": [5.6973677, 45.1518856]  
  },  
  "properties": {  
    "gid": 1,  
    "color" : "green"  
  }  
}  
(1 row)
```

# **EASY TESTING & DEPLOY**



# APT

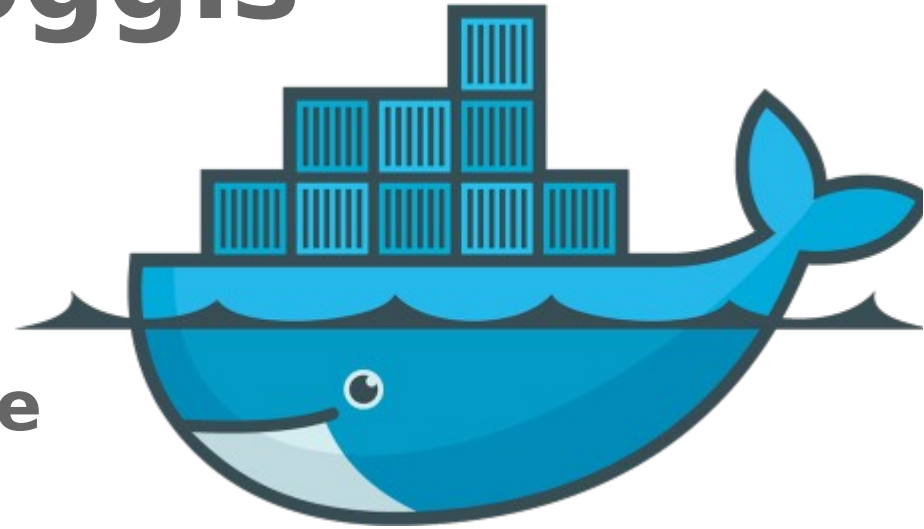


debian

- › **[apt.postgresql.org](http://apt.postgresql.org)**
- › **Supported PostgreSQL versions**
- › **For various apt-based distributions**
- › **+ extensions**
- › **Maintained & Trusted**
- › **PostGIS is in !**
- › **Use it !**



# docker-pggis



- › **Docker container**
- › **Based on Phusion Baseimage**
- › **All GIS-related PG packages**
  - › **PG 9.4**
  - › **PostGIS (+SFCGAL), PgRouting**
  - › **PointCloud, PDAL**

## › **Oneliner run**

```
docker run --rm -P --name pggis_test oslandia/pggis /sbin/my_init
```

**<https://github.com/vpicavet/docker-pggis>**

**More to  
come...**



# ORACLE SPATIAL FDW



# Oracle Spatial FDW

- › **Oracle FDW**
- › **Transparent access to Oracle tables**
- › **No Oracle spatial support**
  - › **WKT as intermediary format**
  - › **Complex and inefficient**
- › **Native support**
- › **SDO → PostGIS**
- › **Replication, cross-db queries...**
- › **Still WIP**

**WIP**

**[https://github.com/laurenz/oracle\\_fdw/pull/7](https://github.com/laurenz/oracle_fdw/pull/7)**



# Oracle Spatial FDW

```
-- enable extension
create extension oracle_fdw;

-- Oracle server parameters and user mapping
create server orcl foreign data wrapper oracle_fdw options (dbserver '//10.0.0.1/FG');
grant usage on foreign server orcl to pggis;
create user mapping for pggis server orcl options (user 'PGGIS', password 'PGGISORAPASS');

-- create foreign table
CREATE FOREIGN TABLE ora_ways (
    way_id bigint,
    description varchar,
    geom geometry(LineString, 2154)
) SERVER orcl OPTIONS(schema 'MYUSER', table 'ways');

-- Use the foreign table just as another table
select * from ora_ways where st_length(geom) < 0.00003;
```



# **BETTER SPATIAL INDEXES**



# SP-Gist

- › « SP » = SPatial
- › New PG index type
- › Faster to read
- › 3x faster to build
- › Good fit for spatial data
- › GSoC 2014
- › Commitable code ?



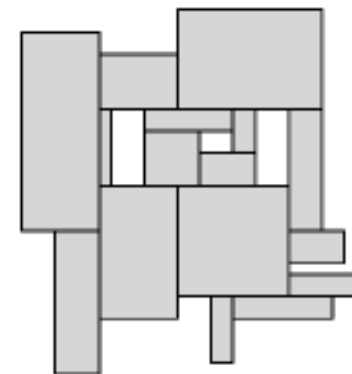
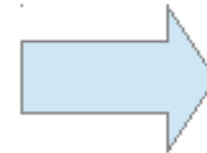
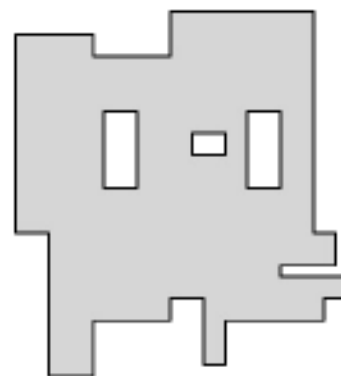
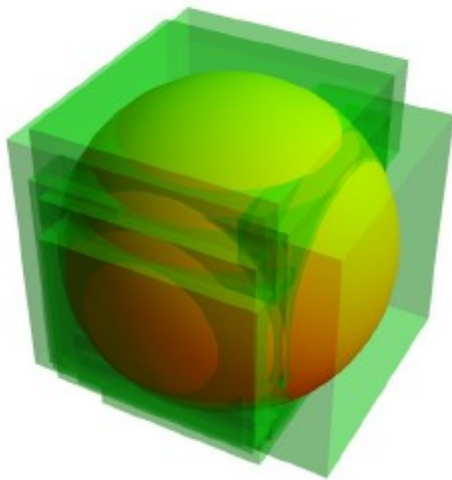


# VODKA



- › **Korotkov, Bartunov, Sigaev**
- › **create index .. using vodka**
- › **Derivation of JSONB indexing**
- › **R-Tree based on GiST as entry tree**
- › **Use multiple boxes per polygons**

**WIP**



# Thank you



**vincent.picavet@oslandia.com**

OSLANDIA

**@vpicavet**

**www.oslandia.com**

