



# PostGIS 2.0 ...

... and beyond

**Vincent Picavet – Oslandia**

[https://github.com/Oslandia/presentations/tree/master/fosdem\\_2013](https://github.com/Oslandia/presentations/tree/master/fosdem_2013)

# PostGIS 2.0

PostGIS 2.0.0 : April 3, 2012

After 26 months !

Major version

Breaks compatibility

Loads of new features

Performance  
improvement





# What's new ?



# Internals

**New serialization format**

**New geometry types (3D)**

**Fix 2D only bounding boxes**

**Fix bytes alignment**

**New parsers**

**WKB**

**WKT**

# Features

Management functions

ISO SQL/MM compliancy

New functions for analysis

Topology (SQL/MM)

Real 3D storage

Raster / geometry functions

KNN indexed search

TIGER (geocoder / reverse...)

# Install

Easier installation (PG  $\geq$  9.1)

```
CREATE EXTENSION postgis ;
```

```
CREATE EXTENSION postgis_topology ;
```



# Manage

geometry\_columns → view

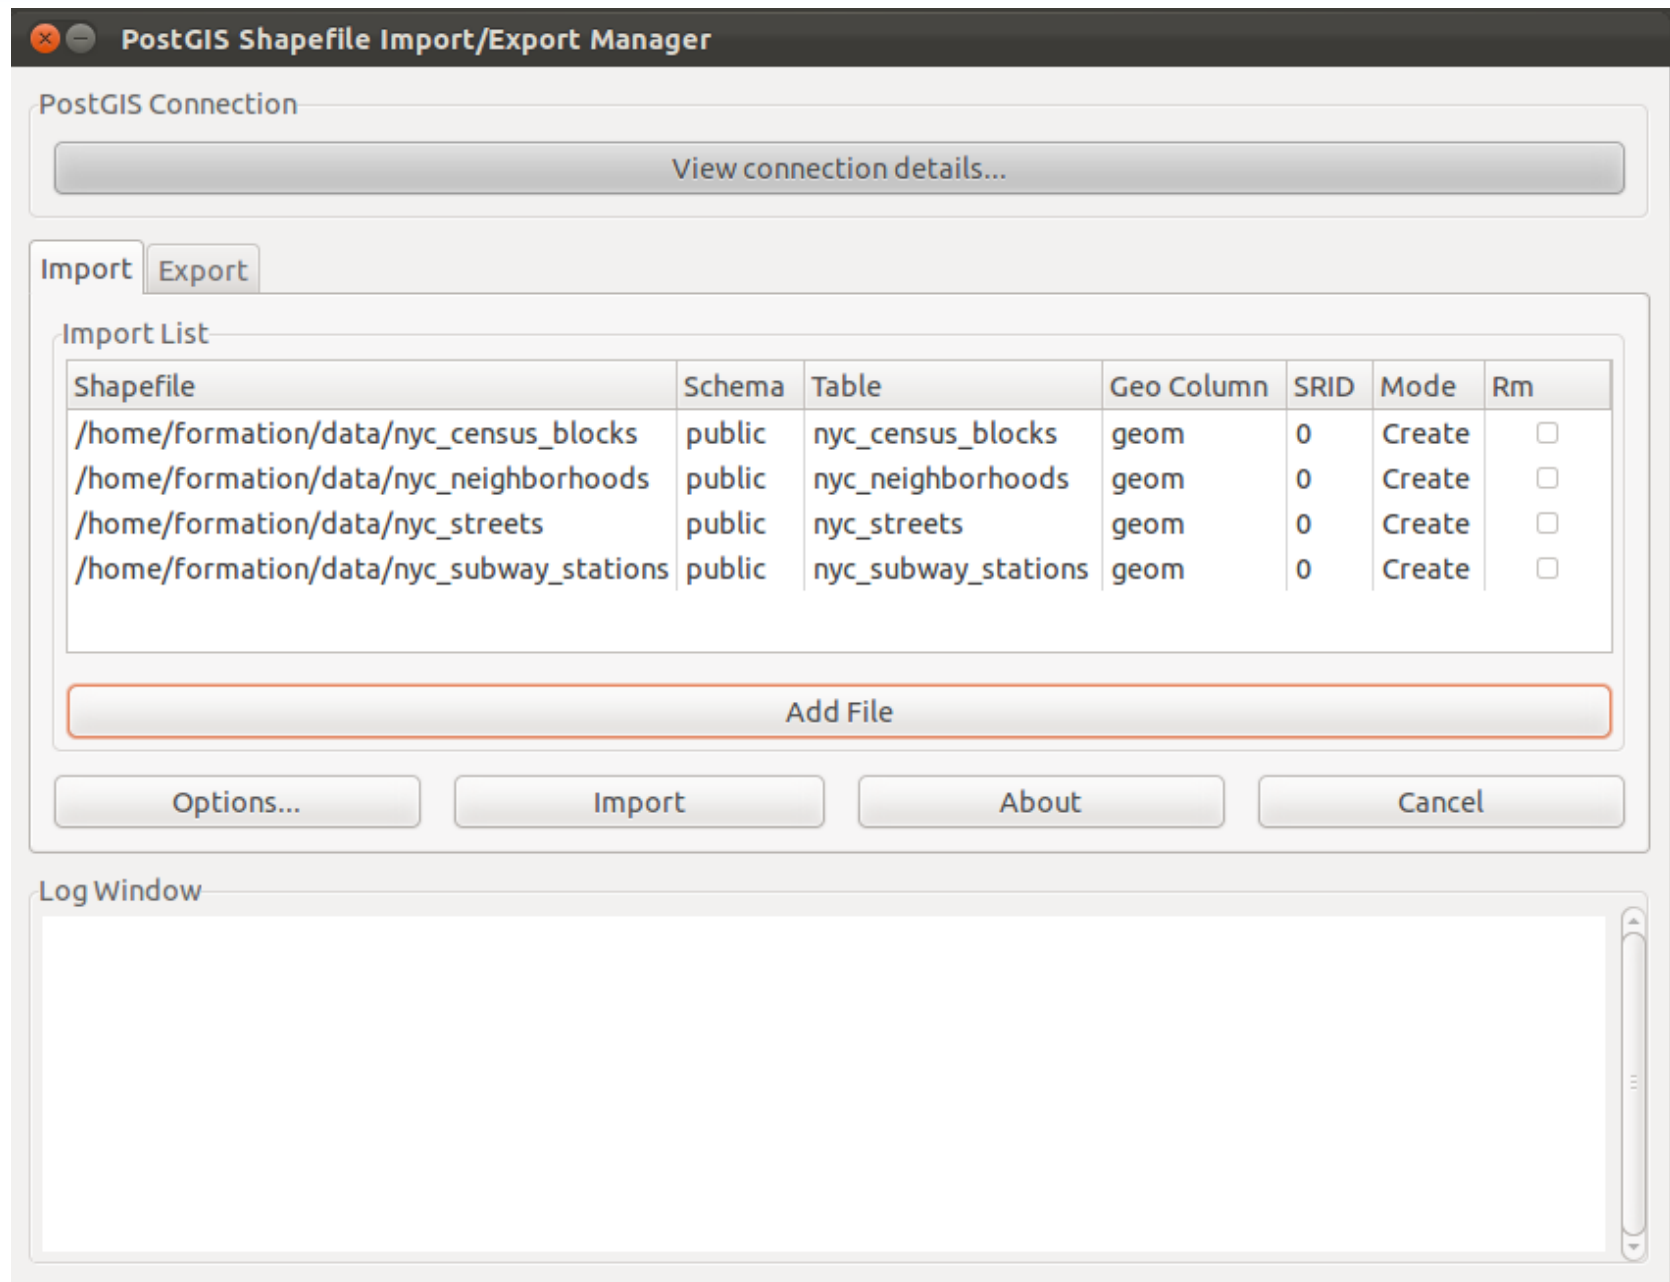
Typmod usage

Old way still available

```
CREATE TABLE buildings (  
    gid SERIAL PRIMARY KEY  
    , geom geometry(MultiPolygon, 26986)  
);
```

```
alter table buildings  
    alter column geom  
        type geometry(MultiPolygon, 2154)  
        using st_setsrid(geom, 2154);
```

# Load





# Functions

ST\_ConcaveHull

ST\_Snap

ST\_Split

ST\_MakeValid

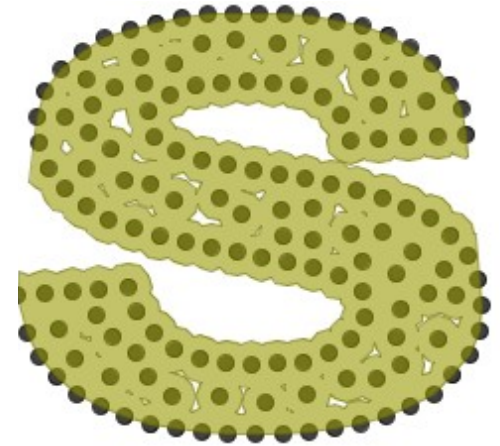
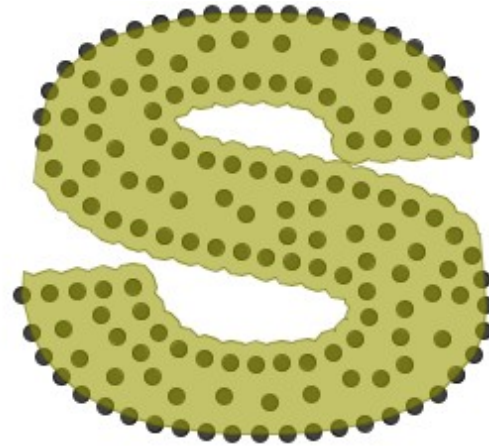
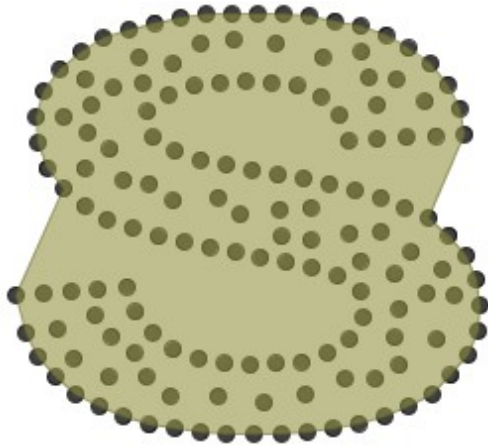
ST\_IsValidDetail

ST\_OffsetCurve

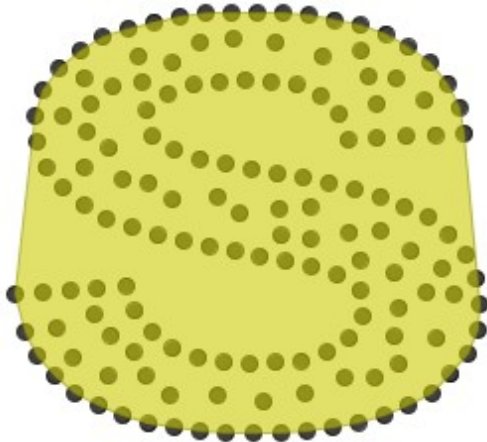
...



# hulls and curves



Concave hulls with various settings



Convex hull



Offset curves

# cleaning data

Before : ST\_Buffer(the\_geom, 0)

After :

ST\_MakeValid()

ST\_RemoveRepeatedPoints()

ST\_IsValidReason()

ST\_IsValidDetail()

```
SELECT ST_IsValid(geom),ST_IsValidReason(geom) FROM
(SELECT ST_GeomFromText('POLYGON ((0 0, 0 10, 10 10, 10 0, 0 0),(20 20, 20 30, 30 30, 30 20, 20 20))') as geom) as foo;
st_isvalid|          st_isvalidreason
-----+-----
```

```
f          | Hole lies outside shell at or near point (20.0, 20.0, NaN)
```

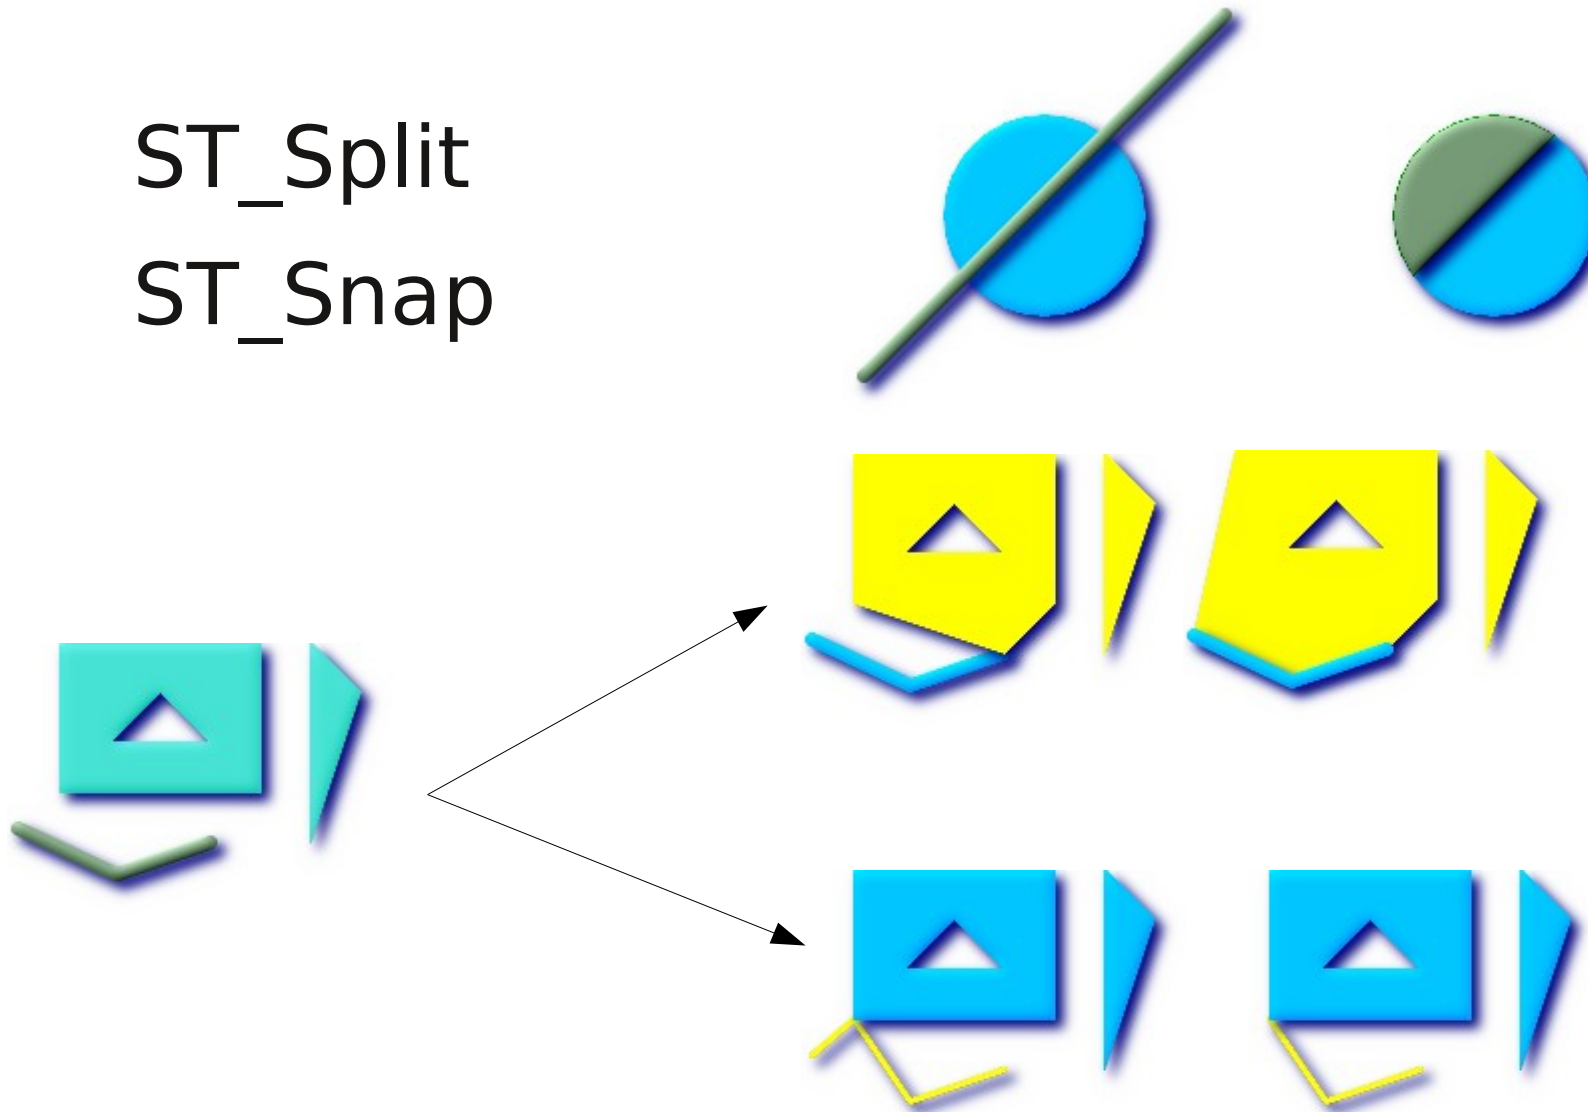
```
SELECT * FROM ST_IsValidDetail('LINESTRING(...)');
```

gid	reason	location
5330	Self-intersection	POINT(32 5)

# Splitting and snapping

ST\_Split

ST\_Snap



# Real 3D

«real» 3D inside PostGIS

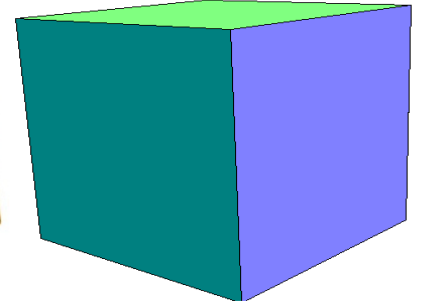
ISO and OGC standards

ISO 19125, SQL/MM, SFS 1.2.0

First step of implementation

New data types & functions

... Much more coming soon ...



New types :

TRIANGLE, POLYHEDRALSURFACE, TIN

New functions :

ST\_3DDistance, ST\_3DIntersects,  
ST\_3DDWithin, ST\_3DClosestPoint...

Input/Output : ST\_AsGML, ST\_AsX3D...

New operators

&&&

Spatial index : nd-indexes

# TIGER

Geocoder

Reverse geocoder

TIGER to PostGIS topology loader

Updated loader





# Topology



*Beware of the spaghetti monster !*



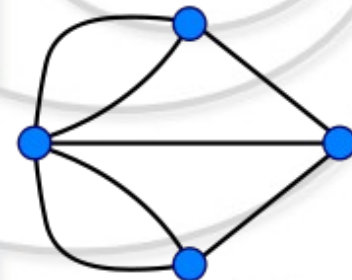
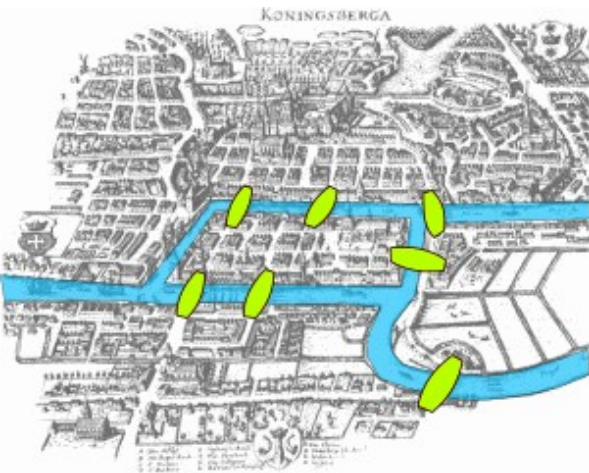
# Topology - Graphs

Explicit relations between objects

Graph representation

Various models

OGC : Node / edge / face



# Topology

Node/Edge/Face model

TopoGeometry Datatype

Use schemas

«topology» for functions and others

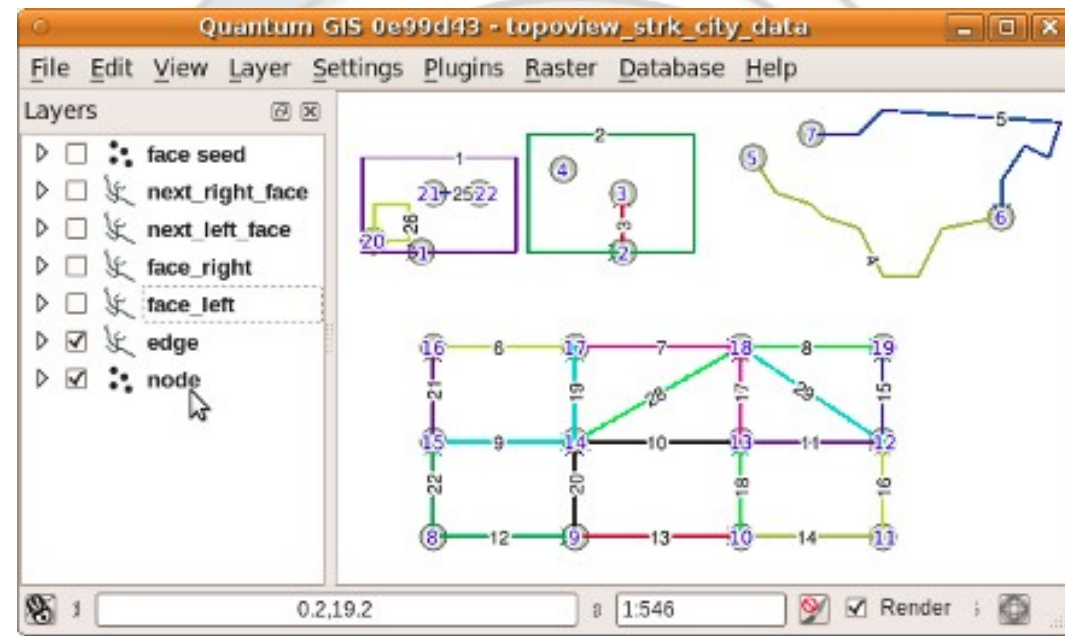
Each topology in its own schema

Full SQL/MM topology support

Integrated in 2.0

Sandro Santilli

Toscane Region





# Topology use case

Table name : **tr**



Fichier Éditer Vue Couche Préférences Extension Vecteur Base de donnée Raster Aide

Couches

- ☐ recursive\_upstream\_topo
- ☐ recursive\_upstream
- ☐ shortest\_path\_topology
- ☐ shortest\_path\_pgrouting
- ☒ hydro network
- ☒ background

Attribute table - hydro network :: 0 / 18936 feature(s) selected

	gid	source	target	hname	cost
0	17681	3042	3041	ruisseau de...	13.1468627...
1	50006	4363	4376	ruisseau de...	154.831357...
2	107308	4427	4443	ruisseau la ...	70.4784694...
3	110767	4810	4816	ruisseau le ...	426.452159...
4	8923	4892	4827	ruisseau de...	1648.21133...
5	109594	5158	5264	rivière la di...	946.014083...
6	45039	5407	5429	NULL	114.028638...
7	105937	5480	5594	ruisseau le ...	824.626701...
8	104620	5481	5518	ruisseau la ...	243.004034...

☒ Contrôle de l'ordre de rendu des couches

```
-- Create a topology
SELECT topology.CreateTopology('hydro', 2154);
-- 1

-- we put the postgis topology features for hydro network in another table
CREATE TABLE tr_topo (gid integer);

-- Add a layer
SELECT topology.AddTopoGeometryColumn('hydro', 'public',
    'tr_topo', 'topogeom', 'MULTILINESTRING');
-- 1

-- Populate the layer and the topology from tr geometry features
INSERT into tr_topo (gid, topogeom)
    SELECT gid, topology.toTopoGeom(geom, 'hydro', 1) FROM tr;
```

- [-] Schémas (3)
  - [-] hydro
    - Collationnements (0)
    - Domaines (0)
    - Configurations FTS (0)
    - Dictionnaires FTS (0)
    - Analyseurs FTS (0)
    - Modèles FTS (0)
    - Fonctions (0)
    - Séquences (5)
    - Tables (4)
      - edge\_data
      - face
      - node
      - relation
    - Fonctions trigger (0)
    - Types (0)
    - Vues (1)
      - edge

```
select * from hydro.edge limit 10;
```

neau sortie

ortie de données

Expliquer (Explain)

Messages

Historique

	edge_id integer	start_node integer	end_node integer	next_left_edge integer	next_right_edge integer	left_face integer	right_face integer	geom geometry(LineString)
1	175256	190369	190361	175230	-175243	0	0	01020000206A080
2	167356	183762	181917	166725	167356	0	0	01020000206A080

```
select * from tr_topo limit 10;
```

eau sortie

ortie de données

Expliquer (Explain)

Messages

gid integer	topogeom topology.topogeometry
116768	(1,1,163704,2)
116767	(1,1,163705,2)



create table

rec\_res2 as

with recursive

search\_graph(edge\_id, start\_node, depth, path, length, cycle) as (

select

g.edge\_id, g.start\_node, 1 as depth, ARRAY[g.edge\_id] as path  
, st\_length(g.geom) as length, false as cycle

from

hydro.edge as g

where

edge\_id = 173832

union all

select

g.edge\_id  
, g.start\_node  
, sg.depth + 1 as depth  
, path || g.edge\_id as path  
, sg.length + st\_length(g.geom) as length  
, g.edge\_id = ANY(path) as cycle

from

hydro.edge as g

join

search\_graph as sg

on

sg.start\_node = g.end\_node

where

not cycle

)

select

sg.\*  
, edge.geom as geom

from

search\_graph as sg

join

hydro.edge as edge

on

sg.edge\_id = edge.edge\_id

limit 1000;

Recursive CTE

# 1 : init

```
select
    g.edge_id, g.start_node, 1 as depth, ARRAY[g.edge_id] as path
    , st_length(geom) as length, false as cycle
from
    hydro.edge as g
where
    edge_id = 173832
union all
```



## 2 : recursive part

select

```
g.edge_id  
, g.start_node  
, sg.depth + 1 as depth  
, path || g.edge_id as path  
, sg.length + st_length(g.geom) as length  
, g.edge_id = ANY(path) as cycle
```

Stack the gid to the path  
for this record

Sum up the cost  
( it's the length here)

from

```
hydro.edge as g
```

join

```
search_graph as sg
```

on

```
sg.start_node = g.end_node
```

where

```
not cycle
```

If the record gid is already  
in the path, we have a cycle

Join result set from  
previous iteration  
to connected upstream  
edges

Do not take elements  
which make a cycle



select

sg.\*  
, edge.geom as geom

from

search\_graph as sg

join

hydro.edge as edge

on

sg.edge\_id = edge.edge\_id

limit 1000;

## 3 : Get results

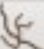





Join CTE results to original table to get geometries

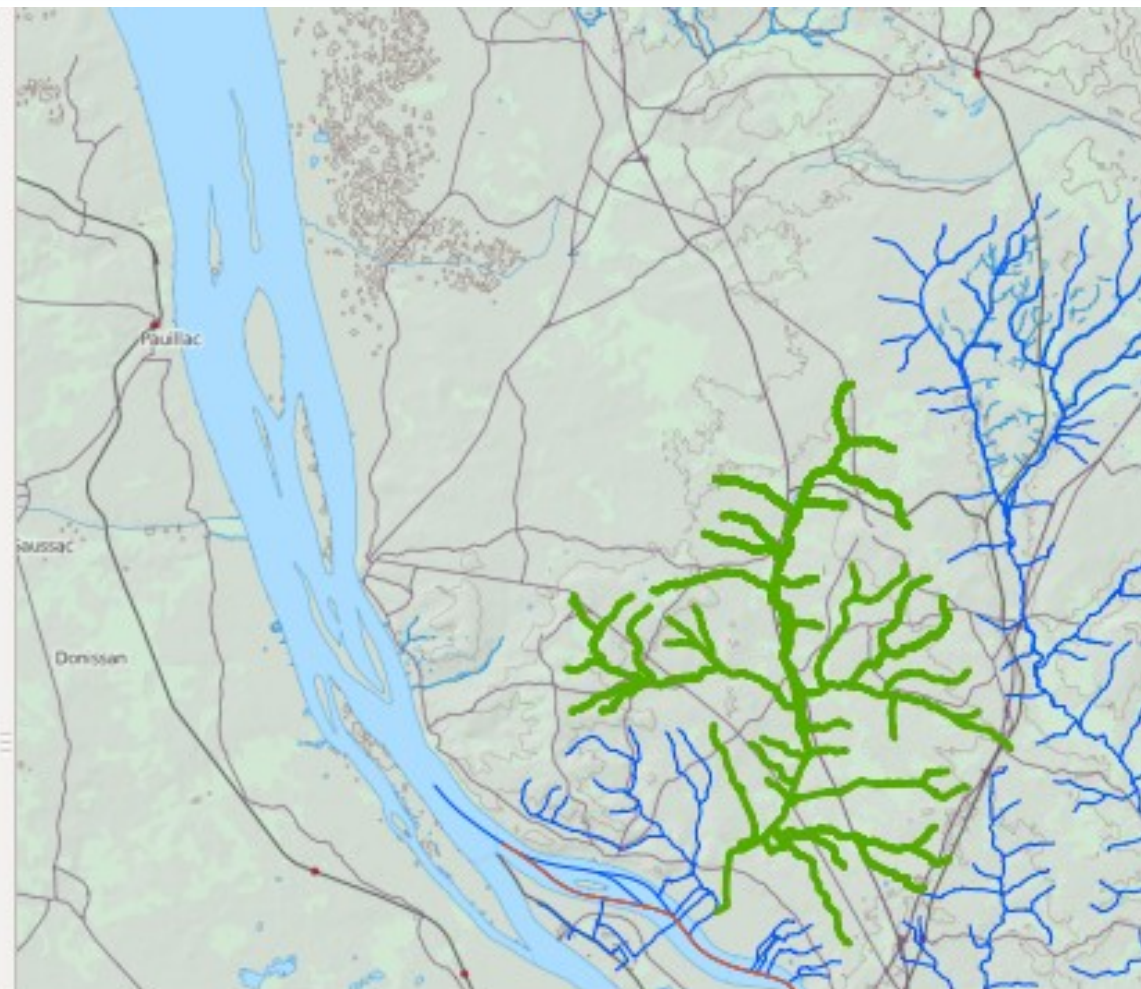
Better limit recursive queries to avoid infinite loops

gid integer	source integer	depth integer	path integer[]	length double precision	cycle boolean	geom geometry(MultiLineString,2154)
31913	20850	1	{31913}	2666.0523017	f	01050000206A080000001000
33855	20735	2	{31913,	3473.3086319	f	01050000206A080000001000
32477	20845	2	{31913,	2725.7640259	f	01050000206A080000001000
33854	19909	3	{31913,	7183.7295195	f	01050000206A080000001000
33853	19909	5	{31913,	6755.7566754	f	01050000206A080000001000



## Couches

- ▶ ☒  recursive\_upstream\_topo
- ▶ ☒  recursive\_upstream
- ▶ ☒  shortest\_path\_topology
- ▶ ☐  shortest\_path\_pgrouting
- ▶ ☒  hydro network
- ▶ ☒  background



Attribute table - recursive\_upstream\_topo :: 0 / 478 feature(s) selected

	edge_id ▲	start_node	depth	path	length	cycle
0	173832	189333	1	{173832}	2666.05230...	f
1	173452	189332	2	{173832,17...	3473.30863...	f



# PostGIS Raster

Raster / vector analysis

New datatype

- Looks like geometry

- But for rasters

Multiresolution, multiband, tile coverage

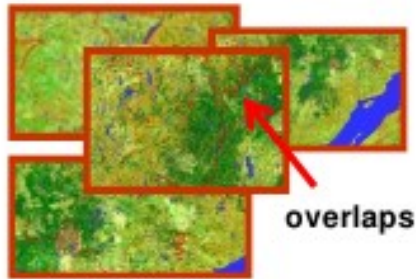
Import/export (GDAL)

Functions

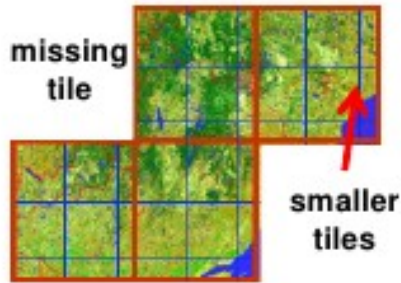
- Statistics, reprojection, edit, compute

- Vector/raster functions

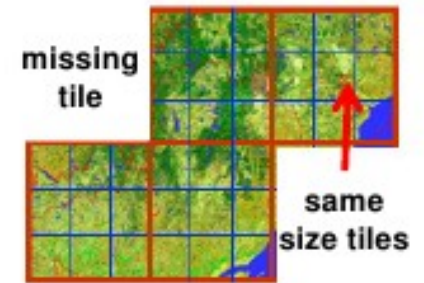
# PostGIS Raster



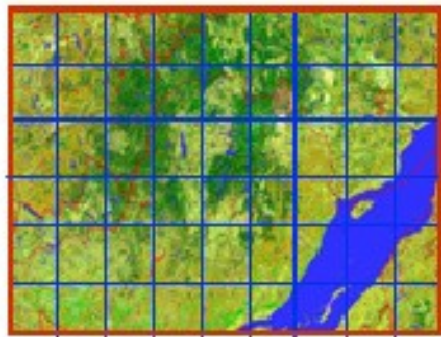
a) warehouse of untiled and unrelated images (4 images)



b) irregularly tiled raster coverage (36 tiles)



c) regularly tiled raster coverage (36 tiles)



d) rectangular regularly tiled raster coverage (54 tiles)

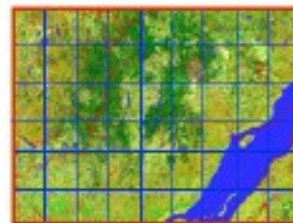


Table 1

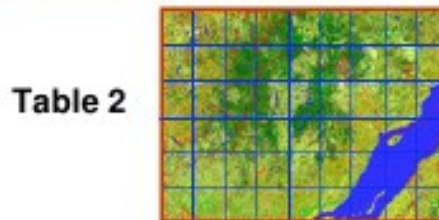
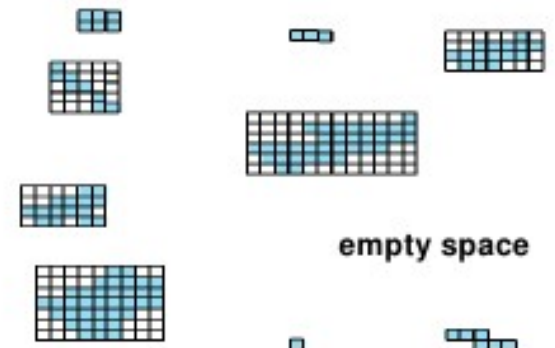


Table 2

e) tiled images (2 tables of 54 tiles)



f) rasterized geometries coverage (9 lines in the table)





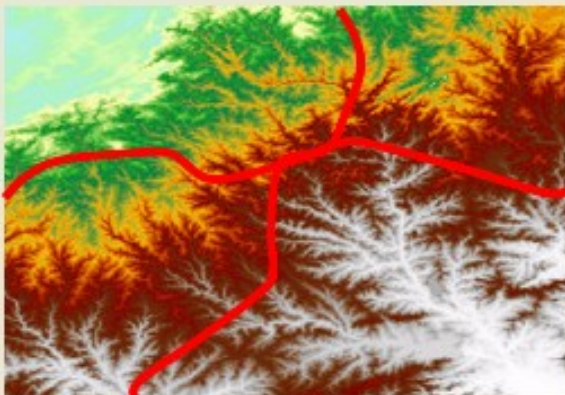
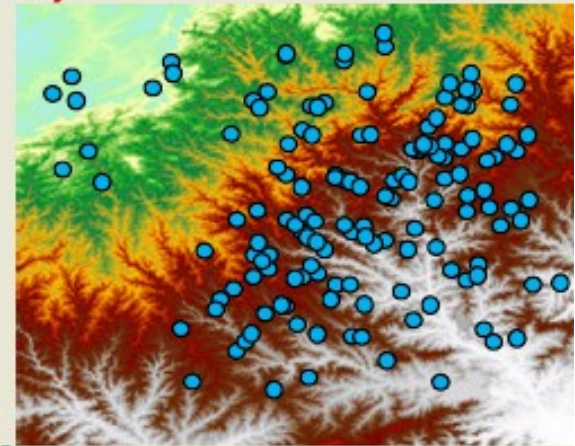
# PostGIS 2.0 : PostGIS Raster

Extract ground elevation values for lidar points...

- `SELECT pointID, ST_Value(rast, geom) elevation`  
`FROM lidar, srtm WHERE ST_Intersects(geom, rast)`

Intersect a road network to extract elevation values for each road segment

- `SELECT roadID,`  
`(ST_Intersection(geom, rast)).geom road,`  
`(ST_Intersection(geom, rast)).val elevation`  
`FROM roadNetwork, srtm WHERE ST_Intersects(geom, rast)`



# PostGIS 2.0 : nearest neighbours

KNN-GIST search in PostgreSQL 9.1

Use indexes !

Spatial nearest neighbors

```
SELECT name, gid FROM geonames
ORDER BY
    geom <-> st_setsrid(st_makepoint(-90,40),4326)
LIMIT 10;
```

Distance operator

<-> or <#> : center or bbox

Need to refine for non-point geometries

# PostGIS 2.1 ?



# PostGIS 2.1

Currently under work :

Arc-geometry distance

Distance with cached tree

R-Tree index improvement (pick-split)

SP-Gist Index

New in PG 9.1

Up to 3x faster construction

Faster to read

# PostGIS 2.1

Under work, raster part :

ST\_MapAlgebra with n rasters

Complete St\_SetValues

Bands to array function

Better « isnodata » management

Raster tiling

(proprietary) support in FME ETL



And more...

Topology improvement

Tiger geocoder as PG extension

PgRouting as PG extension

+ windows support

# PostGIS 2.1

And even more...

Development platform improvement

Meet Debbie and Winnie !

## Jenkins

search

log in | sign up

Jenkins

People

Build History

Build Queue

No builds in the queue.

Build Executor Status

#	Status
1	Idle
2	Idle
3	Idle
4	Idle

AllGDALGEOSPostGISPostgreSQL

S	W	Name ↓	Last Success	Last Failure	Last Duration
		<a href="#">GDAL PostGIS Regress</a>	1 day 17 hr ( <a href="#">#113</a> )	1 mo 1 day ( <a href="#">#2</a> )	5 min 2 sec
		<a href="#">GDAL Regress</a>	1 day 17 hr ( <a href="#">#133</a> )	5 days 15 hr ( <a href="#">#127</a> )	2 min 19 sec
		<a href="#">GDAL Trunk</a>	1 day 17 hr ( <a href="#">#159</a> )	20 days ( <a href="#">#89</a> )	17 min
		<a href="#">GEOS Trunk</a>	1 mo 7 days ( <a href="#">#13</a> )	N/A	2 min 28 sec
		<a href="#">PG Version</a>	11 days ( <a href="#">#9</a> )	1 mo 13 days ( <a href="#">#6</a> )	4 min 7 sec
		<a href="#">PG Version Dev</a>	2 days 10 hr ( <a href="#">#11</a> )	1 mo 1 day ( <a href="#">#4</a> )	4 min 40 sec
		<a href="#">PostGIS 2.0</a>	16 hr ( <a href="#">#34</a> )	4 days 20 hr ( <a href="#">#29</a> )	9 min 24 sec
		<a href="#">PostGIS 2.0 docs</a>	16 hr ( <a href="#">#48</a> )	7 days 3 hr ( <a href="#">#32</a> )	10 min
		<a href="#">PostGIS 2.1</a>	11 hr ( <a href="#">#230</a> )	2 days 21 hr ( <a href="#">#220</a> )	14 min
		<a href="#">PostGIS 2.1 docs</a>	11 hr ( <a href="#">#148</a> )	16 hr ( <a href="#">#145</a> )	10 min
		<a href="#">PostGIS 2.1 doxygen</a>	20 hr ( <a href="#">#45</a> )	8 hr 18 min ( <a href="#">#47</a> )	23 min

# PostGIS 2.1

```
[00:40] <debbie> Project PostGIS_2.0 build #34: SUCCESS in 9 min 21 sec: http://debbie.postgis.net:8080/job/PostGIS\_2.0/build/34
[00:40] <debbie> Paul Ramsey: (#2026) fix performance regression in geography distance calculation
[00:40] <sigq> Title: PostGIS_2.0 #34 [Jenkins] (at debbie.postgis.net:8080)
[00:44] <pramsey> done!
[00:47] <debbie> Project PostGIS_2.1 build #227: SUCCESS in 26 min: http://debbie.postgis.net:8080/job/PostGIS\_2.1/build/227
[00:47] <debbie> Paul Ramsey: (#2026) fix performance regression in geography distance calculation
[00:47] <sigq> Title: PostGIS_2.1 #227 [Jenkins] (at debbie.postgis.net:8080)
[00:51] --> tomkralidis a rejoint ce canal (~tomkralid@CPE0013ce450e14-CM001692413c80.cpe.net.cable.rogers.ca)
[00:51] <-- tomkralidis a quitté ce serveur (Changing host).
[00:51] --> tomkralidis a rejoint ce canal (~tomkralid@osgeo/member/tomkralidis).
[00:51] <debbie> Project PostGIS_2.1_docs build #145: FAILURE in 4 min 0 sec: http://debbie.postgis.net:8080/job/PostGIS\_2.1\_docs/build/145
[00:51] <debbie> * Bborie Park: Added news and docs for ST_Tile(raster). Additional regression tests for
[00:51] <debbie> one additional variant of ST_Tile(raster)
[00:51] <sigq> Title: PostGIS_2.1_docs #145 [Jenkins] (at debbie.postgis.net:8080)
[00:51] <debbie> * Bborie Park: Added ST_Tile() and regression tests. The circle is complete.
[00:51] <debbie> * Bborie Park: Added rt_band_get_pixel_line() and regression tests
[00:51] <debbie> * Paul Ramsey: (#2063) fix the vertex-crossing logic in the circular tree code to use the new edge
[00:51] <debbie> * Paul Ramsey: (#2026) fix performance regression in geography distance calculation
[00:52] --> tbowden a rejoint ce canal (~tim@124-148-118-242.dyn.iinet.net.au).
[00:54] <-- epifanio a quitté ce serveur (Read error: Operation timed out).
[00:58] <winnie> Project PostGIS 2.1 mingW regress build #456: STILL FAILING in 2 min 30 sec: http://winnie.postgis.net:1500/job/PostGIS\_2.1\_mingW\_regress/build/456
[00:58] <sigq> Title: PostGIS 2.1 mingW regress #456 [Jenkins] (at winnie.postgis.net:1500)
```

+ Hallie :  
documentation bot ( PostgreSQL FTS & more)

# **And beyond ?**



# PostGIS 3.0 ?

Paris codesprint and barcamp May 2012

Find directions for future

- Git, build system (mainly Windows)

- Geometry backend (GEOS vs BGL vs ?)

- Raster improvement

- 3D topology & processing (CGAL?)

- Point clouds

- Performance, performance, performance

**Let's go 3D !**



# PostGIS 3D

2.5D already in  
3D storage is in

We want analysis !

ST\_3Dintersects

ST\_3Dintersection

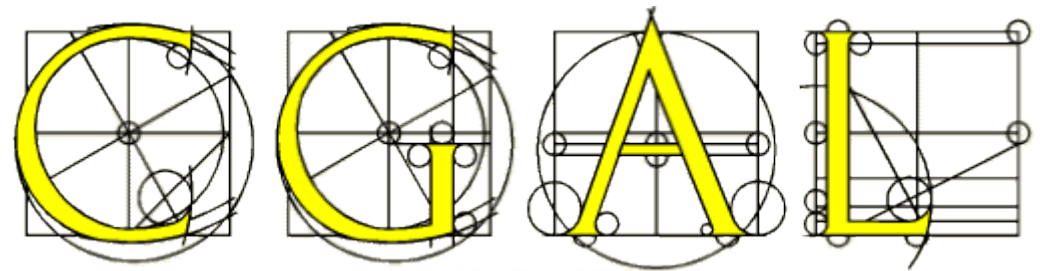
ST\_extrude (2D → 3D)

ST\_3Dconvexhull

ST\_DelaunayTriangles...







2D & 3D geometric computation

C++

**Exact** computation

Efficient, generic, extensible...



**...and now GPL !**



# PostGIS 3D

Some European funding  
Cooperation with IGN & others

e-PLU

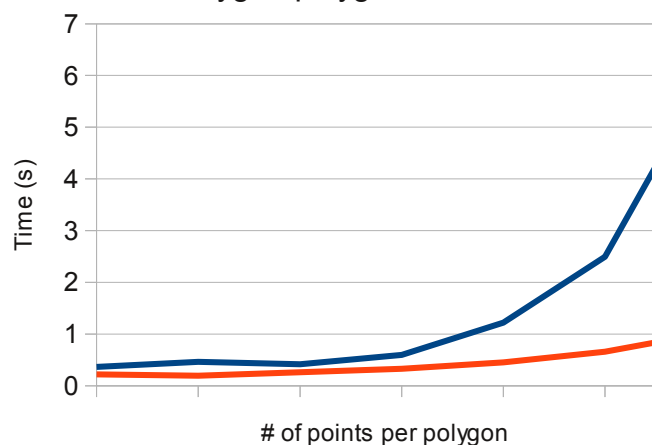
Creation of SFCGAL framework (OGC-SF)  
Use it in PostGIS  
Compare with GEOS (for 2D)



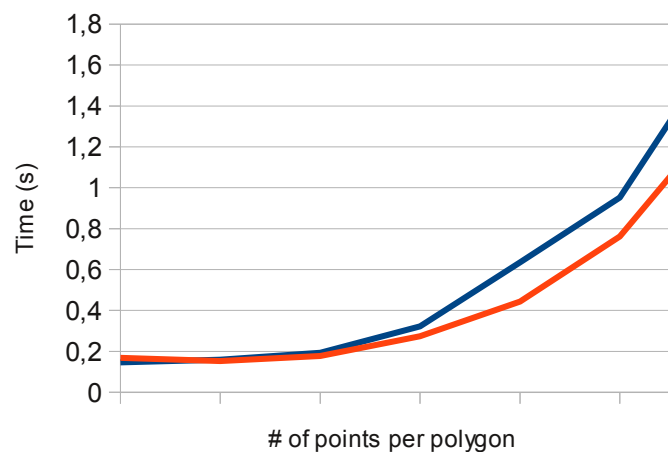
# PostGIS 3D

Most operations faster  
Some operations to optimize

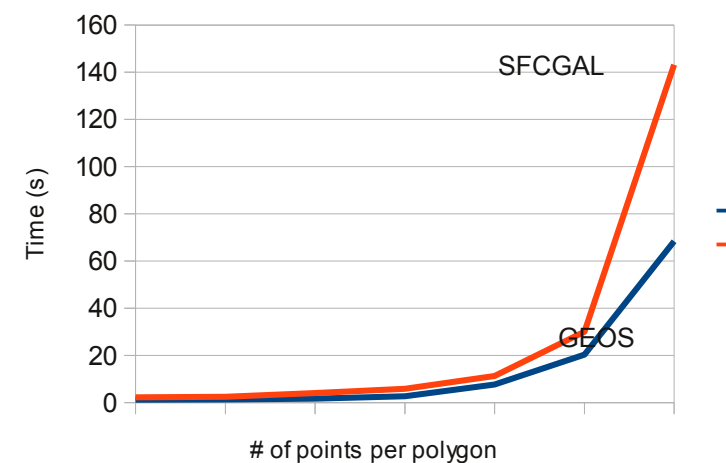
Polygon/polygon intersection



Convex hull



Linestring/polygon intersection



— GEOS  
— SFCGAL

## Issues

CGAL : exact predicates & constructions

PostGIS / GEOS : snap bounding (2D)

```
SELECT
  ST_Intersects(
    ST_Intersection(
      'LINESTRING(0 0,2 1)::geometry',
      'LINESTRING(1 0,0 1)::geometry'),
    'LINESTRING(0 0,2 1)::geometry');
st_intersects
-----
f
(1 row)
```

ry reference

See Hugo's talk on GitHub

# (very) Recent progress

Quantum GIS client ( Globe )

Some analysis functions

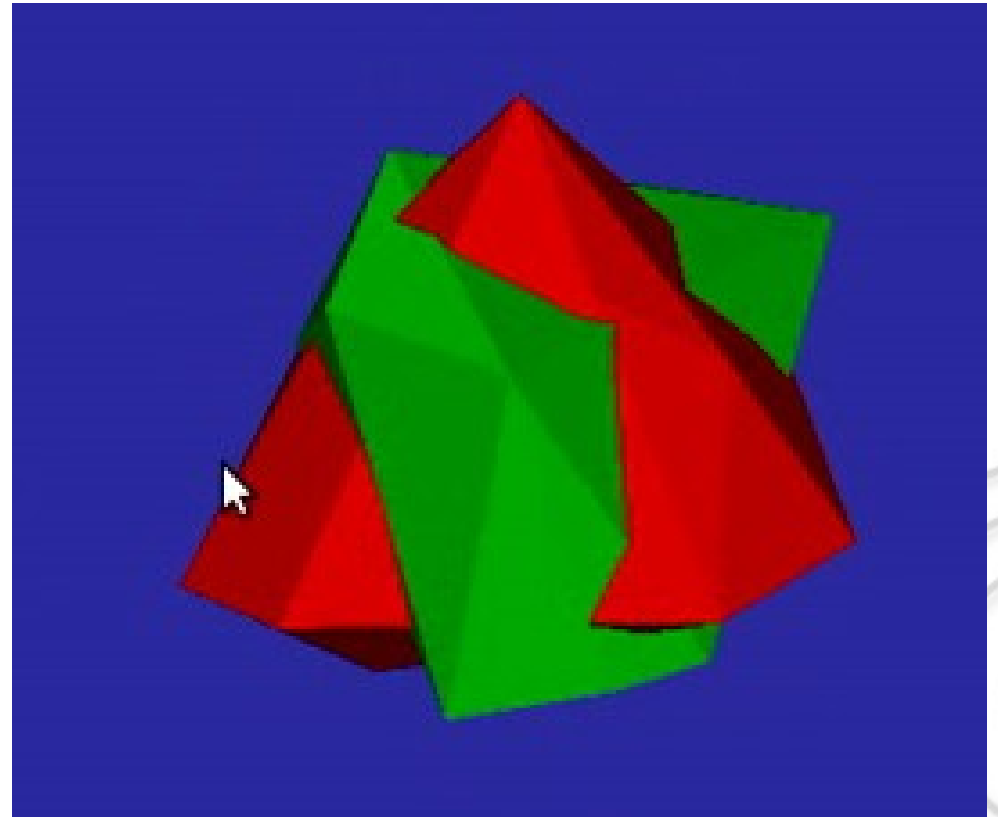
ST\_Extrude

ST\_3DConvexhull

ST\_3DIntersection

Surfaces

Solids





# Want to see ?



# 3D Next steps

Debug

More features from CGAL

Better QGIS support

CityGML & Collada loaders / exporters

PostGIS core integration

To be discussed at Boston codesprint 2013

TinyOWS for 3D webservices

Find €€€€€ to speed up development

# That's it...

## Questions ?

**[vincent.picavet@oslandia.com](mailto:vincent.picavet@oslandia.com)**

**Twitter : @vpicavet**

**<http://www.oslandia.com>**